

Course Reminders

- Due this Friday (11:59 PM)
 - **PROJECT GROUP SIGNUP**
 - D1
 - A1
- Projects
 - You will be assigned a GitHub repo this weekend - please accept the invitation (it will expire)
 - You will also be assigned a previous project to review (links will be on Canvas)



Australian Bureau of Statistics

1800.0 Australian Marriage Law Postal Survey, 2017

Released on 15 November 2017

Table 5 Participation by Federal Electoral Division(a), Males and Age **Gender apartheid**

Table junk

Yeah NA		18-19 years	20-24 years	25-29 years	30-34 years	35-39 years	40-44 years	45-49 years	50-54 years	55-59 years	60-64 years
Lingian (c)	Total participants	292	1,058	1,465	1,653	1,515	1,516	1,710	1,730	1,753	1,574
	Eligible participants	572	2,910	3,789	3,996	3,607	3,506	3,645	3,331	2,960	2,456
	Participation rate (%)	51.0	36.4	38.7	41.4	42.0	43.2	46.9	51.9	59.2	64.1
Primary keynotes											
Merged cells	Total participants	442	1,461	2,066	2,357	2,188	2,057	2,224	2,108	2,134	1,772
	Solomon	750	2,991	3,994	4,155	3,634	3,398	3,427	3,066	2,931	2,355
	Eligible participants	58.9	48.8	51.7	56.7	60.2	60.5	64.9	68.8	72.8	75.2
Northern Territory (Total)	Total participants	734	2,519	3,531	4,010	3,703	3,573	3,934	3,838	3,887	3,346
	Eligible participants	1,322	5,901	7,783	8,151	7,241	6,904	7,072	6,397	5,891	4,811
	Participation rate (%)	55.5	42.7	45.4	49.2	51.1	51.8	55.6	60.0	66.0	69.5
Australian Capital Territory Divisions	Covariate as Subheading	Summary of data inside data									
	Canberra(d)	Total participants	1,764	4,789	4,817	4,973	4,626	4,453	5,074	4,826	5,169
		Eligible participants	2,260	6,471	6,448	6,509	5,983	5,805	6,302	5,902	6,044
Fenner(e)	Participation rate (%)	78.1	74.0	74.7	76.4	77.3	76.7	80.5	81.8	85.5	86.9
	Total participants	1,477	4,687	5,178	5,786	6,025	5,463	5,191	4,208	3,948	3,465
	Eligible participants	1,904	6,354	7,121	7,822	7,960	7,155	6,480	5,206	4,692	3,945
Australian Capital Territory (Total)	Participation rate (%)	77.6	73.8	72.7	74.0	75.7	76.4	80.1	80.8	84.1	87.8
NA Yeah											
Total participants	3,241	9,476	9,595	10,759	10,031	9,910	10,263	9,034	9,117	7,859	
Australia	Eligible participants	4,164	12,825	13,569	14,331	13,943	12,960	12,782	11,108	10,736	9,002
	Participation rate (%)	77.8	73.9	73.7	75.1	76.4	76.5	80.3	81.3	84.9	87.3
	Total participants	151,297	438,166	441,658	460,548	462,206	479,360	524,620	517,693	543,449	506,799
Total	Eligible participants	201,439	635,909	646,916	665,250	656,446	660,841	693,850	659,150	664,720	597,386
	Participation rate (%)	75.1	68.9	68.3	69.2	70.4	72.5	75.6	78.5	81.8	84.8

(a) The Federal Electoral Divisions are current as at 24 August 2017.

(b) Includes those whose age is unknown

(c) Includes Christmas Island and the Cocos (Keeling) Islands

(d) Includes Norfolk Island.

(e) Includes Norfolk Island

Return of the table junk

untidy data

Australian Bureau of Statistics												
1800.0 Australian Marriage Law Postal Survey, 2017												
Released 15 November 2017												
Table junk												
Yeah NA												
Lingua (c)	Total participants	292	1,058	1,460	1,653	1,515	1,516	1,710	1,730	1,753	1,574	1,574
Lingua (c)	Eligible participants	572	2,910	3,049	3,996	3,607	3,506	3,645	3,331	2,960	2,456	2,456
Primary keynote	Participation rate (%)	51.0	36.4	38.7	41.4	42.0	43.2	46.8	51.9	58.2	64.1	64.1
Merged cells	Comma on											
Solomon	Total participants	442	1,461	2,068	2,357	2,188	2,057	2,224	2,108	2,134	1,772	1,772
Solomon	Eligible participants	750	2,991	3,994	4,155	3,634	3,398	3,427	3,066	2,931	2,355	2,355
Solomon	Participation rate (%)	56.9	48.8	51.7	56.7	60.2	60.5	64.9	68.8	72.8	75.2	75.2
Northern Territory	Total participants	734	2,519	3,531	4,010	3,703	3,573	3,934	3,838	3,887	3,346	3,346
(Total)	Eligible participants	1,322	5,901	7,783	8,151	7,241	6,904	7,072	6,397	5,891	4,811	4,811
(Total)	Participation rate (%)	55.5	42.7	45.4	49.2	51.1	51.8	55.6	60.0	66.0	69.5	69.5
Australian Capital Territory Divisions	Covariate as Subheading											
Canberra(d)	Total participants	1,764	4,789	4,817	4,973	4,626	4,453	5,074	4,826	5,169	4,394	4,394
Canberra(d)	Eligible participants	2,260	6,471	6,448	6,509	5,983	5,805	6,302	5,902	6,044	5,057	5,057
Canberra(d)	Participation rate (%)	78.1	74.0	74.7	76.4	77.3	76.7	80.5	81.8	85.5	86.9	86.9
Fenner(e)	Total participants	1,477	4,687	5,178	5,786	6,025	5,463	5,191	4,208	3,948	3,465	3,465
Fenner(e)	Eligible participants	1,904	6,354	7,121	7,822	7,960	7,155	6,480	5,206	4,692	3,945	3,945
Fenner(e)	Participation rate (%)	77.6	73.8	72.7	74.0	75.7	76.4	80.1	80.8	84.1	87.8	87.8
Australia	NA Yeah											
Australian Capital Territory (Total)	Total participants	5,542	9,416	9,895	10,155	10,054	9,219	10,205	9,034	9,141	7,659	7,659
Australian Capital Territory (Total)	Eligible participants	4,164	12,825	13,569	14,331	13,943	12,960	12,782	11,108	10,736	9,002	9,002
Australian Capital Territory (Total)	Participation rate (%)	77.8	73.9	73.7	75.1	76.4	76.5	80.3	81.3	84.9	87.3	87.3
a) The Federal Electoral Divisions are current as at 24 August 2017												
b) Includes those whose age is unknown												
c) Includes Christmas Island and the Cocos (Keeling) Islands												
d) Includes Norfolk Island												
e) Includes Jervis Bay												
Return of the table junk												
MS Excel or Die												

data
wrangling

tidy data

area	gender	age	State	Area (sq km)	Eligible participants	Participation rate (%)	Total participants	Total Participants
Adelaide	Female	18-19 years	SA	76	1341	83.5	1120	1120
Adelaide	Female	20-24 years	SA	76	4620	81.2	3750	3750
Adelaide	Female	25-29 years	SA	76	4897	81.8	4004	4004
Adelaide	Female	30-34 years	SA	76	4784	79.8	3820	3820
Adelaide	Female	35-39 years	SA	76	4319	79	3411	3411
Adelaide	Female	40-44 years	SA	76	4310	80.6	3472	3472
Adelaide	Female	45-49 years	SA	76	4579	81.4	3728	3728
Adelaide	Female	50-54 years	SA	76	4475	84.7	3791	3791
Adelaide	Female	55-59 years	SA	76	4622	87.3	4033	4033
Adelaide	Female	60-64 years	SA	76	4342	89.3	3879	3879
Adelaide	Female	65-69 years	SA	76	3970	90.7	3602	3602
Adelaide	Female	70-74 years	SA	76	3009	90.3	2716	2716
Adelaide	Female	75-79 years	SA	76	2156	88.5	1908	1908
Adelaide	Female	80-84 years	SA	76	1673	85.1	1423	1423

1	area	gender	age	State	Area (sq km)	Eligible participants	Participation rate (%)	Total participants	Total Participants
2	Adelaide	Female	18-19 years	SA	76	1341	83.5	1120	1120
3	Adelaide	Female	20-24 years	SA	76	4620	81.2	3750	3750
4	Adelaide	Female	25-29 years	SA	76	4897	81.8	4004	4004
5	Adelaide	Female	30-34 years	SA	76	4784	79.8	3820	3820
6	Adelaide	Female	35-39 years	SA	76	4319	79	3411	3411
7	Adelaide	Female	40-44 years	SA	76	4310	80.6	3472	3472
8	Adelaide	Female	45-49 years	SA	76	4579	81.4	3728	3728
9	Adelaide	Female	50-54 years	SA	76	4475	84.7	3791	3791
10	Adelaide	Female	55-59 years	SA	76	4622	87.3	4033	4033

A	B	C	D	E	F	G	H	I	J	K	L
 Australian Bureau of Statistics	Table junk										
1800.0 Australian Marriage Law Postal Survey, 2017											
Released on 31 November 2017											
Table 5 Participation by Federal Electoral Division (a) Males and Age Gender apartheid											
Yeah NA	18-19 years	20-24 years	25-29 years	30-34 years	35-39 years	40-44 years	45-49 years	50-54 years	55-59 years	60-64 years	
522	Total participants	292	1,059	3,409	1,053	3,115	1,516	1,710	1,720	1,753	1,574
523	Eligible participants	292	1,059	3,409	1,053	3,107	1,516	1,707	1,717	1,750	1,569
524	Participation rate (%)	51.0	36.4	37.9	51.4	42.0	43.2	46.9	51.9	59.2	44.1
525	Comment on										
526	Merged cells										
527	Solomon										
528	Total participants	443	1,461	3,066	2,367	2,188	2,087	2,234	2,108	2,334	1,771
529	Eligible participants	443	1,461	3,064	2,365	3,634	3,155	3,398	3,427	3,006	2,355
530	Participation rate (%)	58.9	48.8	51.7	56.7	60.2	60.5	64.9	68.8	72.8	75.2
531	Northern Territory (Total)										
532	Total participants	734	2,539	3,551	4,010	3,703	3,573	3,504	3,883	3,887	3,346
533	Eligible participants	734	2,539	3,550	4,010	3,702	3,572	3,503	3,884	3,888	3,423
534	Participation rate (%)	55.5	42.7	45.4	49.2	51.1	51.8	55.6	60.0	68.0	69.5
535	Australian Capital Territory Divisions										
536	Coverlate as Subheading										
537	Summary of territory inside data										
538	Canberra										
539	Total participants	7,745	4,793	4,817	4,973	4,626	4,453	5,074	4,826	5,369	4,294
540	Eligible participants	2,260	4,471	4,648	4,509	5,983	5,805	4,902	5,962	6,044	5,057
541	Participation rate (%)	78.1	74.0	74.7	75.4	77.3	76.7	80.5	81.8	85.5	88.9
542	Fairwater										
543	Total participants	1,477	4,687	5,178	5,786	6,025	5,463	5,191	4,206	3,948	3,465
544	Eligible participants	1,304	4,584	7,212	7,022	7,960	7,155	4,480	5,206	4,602	3,945
545	Participation rate (%)	71.4	72.1	73.0	73.8	75.3	74.8	80.0	81.4	84.1	87.0
546	2										
547	NA Yeah										
548	Total participants	2,591	2,604	2,609	2,601	2,600	2,600	2,600	2,601	2,600	2,600
549	Eligible participants	2,591	2,605	2,609	2,601	2,600	2,600	2,600	2,601	2,600	2,600
550	Participation rate (%)	77.8	73.9	73.7	75.1	78.4	76.5	80.3	81.3	84.9	87.3
551	Australia										
552	Total participants	151,297	438,165	441,168	460,549	462,206	479,360	521,850	517,893	545,149	506,759
553	Eligible participants	201,439	638,909	646,516	665,250	658,446	660,841	681,850	659,150	664,720	597,386
554	Participation rate (%)	75.1	68.9	68.3	69.2	70.4	72.5	75.6	78.5	81.8	84.8

a) The Federal Electoral Divisions are current as at 24 August 2017

b) Includes those whose age is unknown

c) Includes Tasmania and the Cocos (Keeling) Islands

d) Includes Norfolk Island

e) Includes Jervis Bay

Return of the table junk

Survey-data-with-r-5d35cea07962

MS Excel or Die

Tidy Data

1. Each **variable** you measure should be in a single column

	A	B	C	D	E	F	G
1	ID	LastName	FirstName	Sex	City	State	Occupation
2	1004	Smith	Jane	female	Frederick	MD	Welder
3	4587	Nayef	Mohammed	male	Upper Darby	PA	Nurse
4	1727	Doe	Janice	female	San Diego	CA	Doctor
5	6879	Jordan	Alex	male	Birmingham	AL	Teacher

2. Every **observation** of a variable should be in a different row

	A	B	C	D	E	F	G
1	ID	LastName	FirstName	Sex	City	State	Occupation
2	1004	Smith	Jane	female	Frederick	MD	Welder
3	4587	Nayef	Mohammed	male	Upper Darby	PA	Nurse
4	1727	Doe	Janice	female	San Diego	CA	Doctor
5	6879	Jordan	Alex	male	Birmingham	AL	Teacher

3. There should be one table for each type of data

Demographic Survey Data

	A	B	C	D	E	F	G
1	ID	LastName	FirstName	Sex	City	State	Occupation
2		1004	Smith	Jane	female	Frederick	MD
3		4587	Nayef	Mohammed	male	Upper Darby	PA
4		1727	Doe	Janice	female	San Diego	CA
5		6879	Jordan	Alex	male	Birmingham	AL
							Teacher

Doctor's Office Measurements Data

	A	D	E	F	G
1	ID	Height_inches	Weight_lbs	Insulin	Glucose
2		1004	65	180	0.60
3		4587	75	215	1.46
4		1727	62	124	0.72
5		6879	77	160	1.23
					205

4. If you have multiple tables, they should include a column in each *with the same column label* that allows them to be joined or merged

	A	B	C	D	E	F	G
1	ID	LastName	FirstName	Sex	City	State	Occupation
2	1004	Smith	Jane	female	Frederick	MD	Welder
3	4587	Nayef	Mohammed	male	Upper Darby	PA	Nurse
4	1727	Doe	Janice	female	San Diego	CA	Doctor
5	6879	Jordan	Alex	male	Birmingham	AL	Teacher

	A	D	E	F	G
1	ID	Height_inches	Weight_lbs	Insulin	Glucose
2	1004	65	180	0.60	163
3	4587	75	215	1.46	150
4	1727	62	124	0.72	177
5	6879	77	160	1.23	205

Tidy data == rectangular data

A

	A	B	C	D	E
1	id	sex	glucose	insulin	triglyc
2	101	Male	134.1	0.60	273.4
3	102	Female	120.0	1.18	243.6
4	103	Male	124.8	1.23	297.6
5	104	Male	83.1	1.16	142.4
6	105	Male	105.2	0.73	215.7

Tidy Data Benefits

1. consistent data structure
2. foster tool development
3. require only a small set of tools to be learned
4. allow for datasets to be combined

TIDY data is **NOT** the same as **CLEAN** data



A

ID	Last	First	height_m	height_f
1004	Smith	Jane	NA	65
4587	Nayef	Mohammed	72	NA
1727	Doe	Janice	NA	60
6879	Jordan	Alex	55	NA

B

ID	Last	First	height_m	height_f
1004	Smith	Jane		65
4587	Nayef	Mohammed	72	
1727	Doe	Janice		60
6879	Jordan	Alex	55	

C

ID	Last	First	sex	height
1004	Smith	Jane	female	65
4587	Nayef	Mohammed	male	72
1727	Doe	Janice	fem	60
6879	Jordan	Alex	male	55

D

ID	Last	First	sex	height
1004	Smith	Jane	F	65
4587	Nayef	Mohammed	M	72
1727	Doe	Janice	F	60
6879	Jordan	Alex	M	55

Which of these tables stores data best?



A



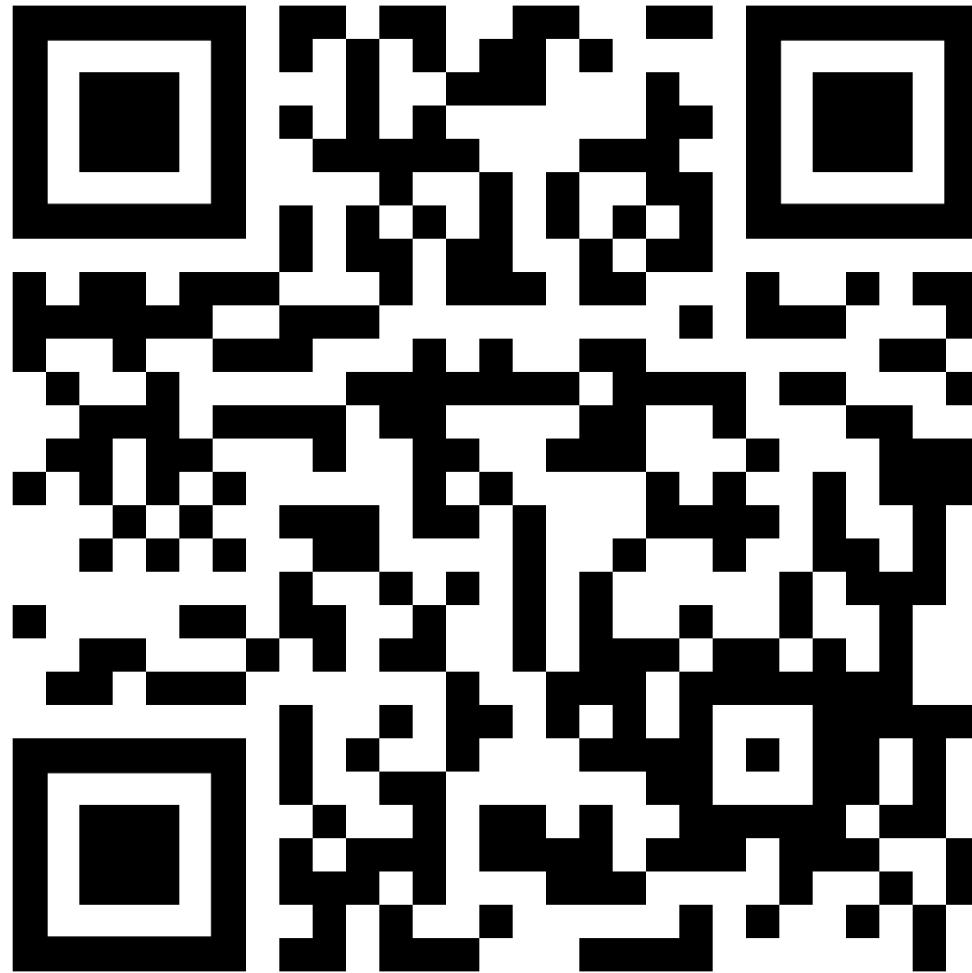
B



C

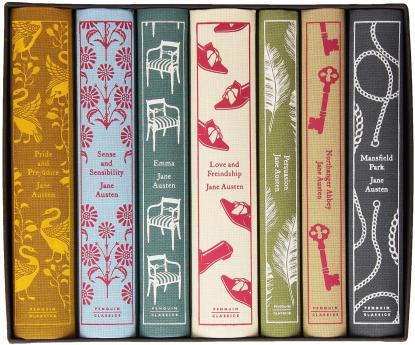


D





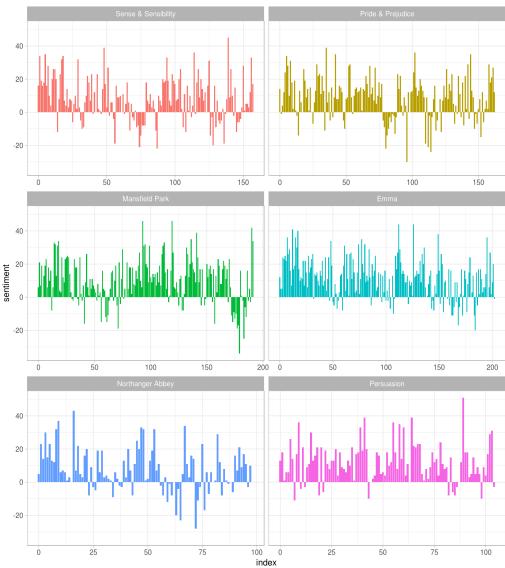
text

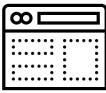


tidy dataset

Word	Novel	Frequency
good	Emma	359
young	Emma	192
friend	Emma	166

results





website

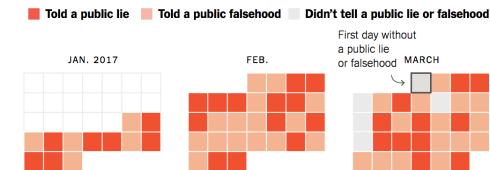
→

tidy dataset

	<u>date</u>	<u>lie</u>	<u>explanation</u>	<u>url</u>
0	Jan. 21, 2017	I wasn't a fan of Iraq. I didn't want to go in.	He was for an invasion before he was against it.	https://www.buzzfeed.com/andrewkaczynski/in-2017-so-far
1	Jan. 21, 2017	A reporter for Time magazine — and I have been... ...	Trump was on the cover 11 times and Nixon appeared 10 times.	http://nation.time.com/2013/11/06/10-things-you-didnt-know-about-nixon/
2	Jan. 23, 2017	Between 3 million and 5 million illegal votes ...	There's no evidence of illegal voting.	https://www.nytimes.com/2017/01/23/us/politics/electoral-college-vote-counts.html
3	Jan. 25, 2017	Now, the audience was the biggest ever. But that...	Official arrival photos show Obama's 2009 inauguration.	https://www.nytimes.com/2017/01/21/us/politics/trumps-inauguration.html
4	Jan. 25, 2017	Take a look at the Pew reports (which show vot...	The report never mentioned voter fraud.	https://www.nytimes.com/2017/01/24/us/politics/trumps-voter-fraud-report.html

—

results



text (lyrics)

ThePudding "I'll be analyzing the repetitiveness of a dataset of 15,000 songs that charted on the Billboard Hot 100 between 1958 and 2017."



AN EXERCISE IN LANGUAGE COMPRESSION

Are Pop Lyrics Getting More Repetitive?

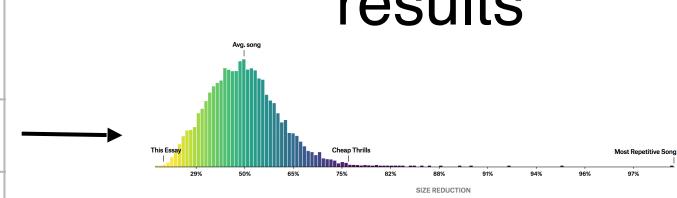
By Colin Morris

<https://pudding.cool/2017/05/song-repetition/>

tidy dataset

song	Artist	Released	Reduction
Cheap Thrills	Sia	2016	76
Around The World	Daft Punk	1997	98
Everybody Dies	J. Cole	2018	27

results



What are these uber-repetitive outliers? *Around The World* by Daft Punk gets reduced a whopping 98%. It goes from 2,610 characters to 61. Small enough to fit in a tweet - twice!

Data Intuition



In today's pattern recognition class my professor talked about PCA, eigenvectors and eigenvalues.

1011



I understood the mathematics of it. If I'm asked to find eigenvalues etc. I'll do it correctly like a machine. But I didn't **understand** it. I didn't get the purpose of it. I didn't get the feel of it.



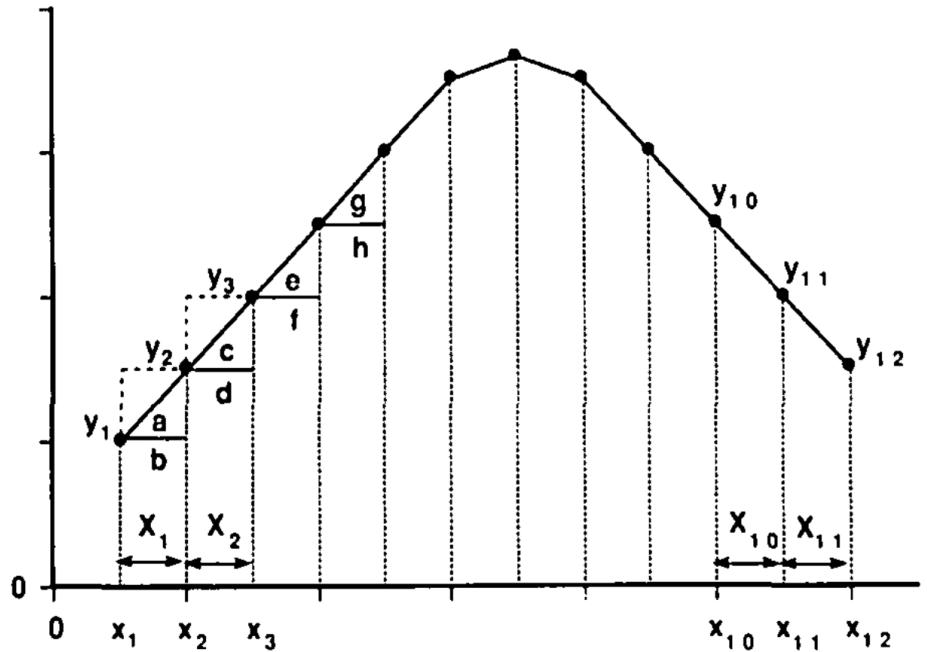
1375



You do not really understand something unless you can explain it to your grandmother. -- Albert Einstein

Well, I can't explain these concepts to a layman or grandma.

1. Why PCA, eigenvectors & eigenvalues? What was the *need* for these concepts?
2. How would you explain these to a layman?



Theory vs. Practice: “Tai’s model”

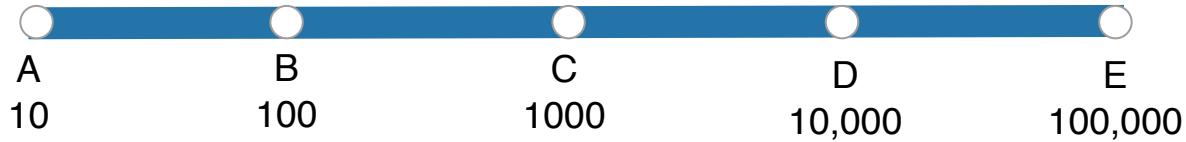
Figure 1—Total area under the curve is the sum of individual areas of triangles *a*, *c*, *e*, and *g* and rectangles *b*, *d*, *f*, and *h*.



Fermi Estimation

<https://forms.gle/C982naWtU9RvHqAb7>

Approximately how many piano tuners do you think there are
in the city of Chicago?







[https://www.youtube.com/watch?
v=0XzyupQXBIs](https://www.youtube.com/watch?v=0XzyupQXBIs)

**Has humanity produced enough
paint to cover the entire land area of
the Earth?**

—Josh (Bolton, MA)



Has humanity produced enough paint to cover the entire land area of the Earth?

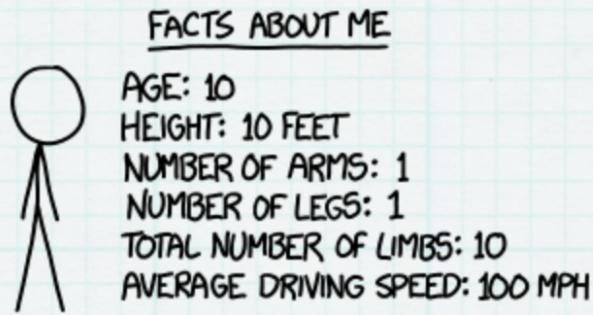




This answer is pretty straightforward. We can look up the size of the world's paint industry, extrapolate backward to figure out the total amount of paint produced. We'd also need to make some assumptions about how we're painting the ground. Note: When we get to the Sahara desert, I recommend not using a brush.



But first, let's think about different ways we might come up with a guess for what the answer will be. In this kind of thinking—often called **Fermi estimation**—all that matters is getting in the right ballpark; that is, the answer should have about the right number of digits. In Fermi estimation, you can round [1] all your answers to the nearest order of magnitude:



Let's suppose that, on average, everyone in the world is responsible for the existence of two rooms, and they're both painted. My living room has about 50 square meters of paintable area, and two of those would be 100 square meters. 7.15 billion people times 100 square meters per person is a little under a trillion square meters —an area smaller than Egypt.

NOT ENOUGH	EXACTLY ENOUGH	MORE THAN ENOUGH
/		

Let's make a wild guess that, on average, one person out of every thousand spends their working life painting things. If I assume it would take me three hours to paint the room I'm in, [2] and 100 billion people have ever lived, and each of them spent 30 years painting things for 8 hours a day, we come up with 150 trillion square meters ... just about exactly the land area of the Earth.

NOT ENOUGH	EXACTLY ENOUGH	MORE THAN ENOUGH
/	/	

How much paint does it take to paint a house? I'm not enough of an adult to have any idea, so let's take another Fermi guess.

Based on my impressions from walking down the aisles, home improvement stores stock about as many light bulbs as cans of paint. A normal house might have about 20 light bulbs, so let's assume a house needs about 20 gallons of paint.^[3] Sure, that sounds about right.

The average US home costs about \$200,000. Assuming each gallon of paint covers about 300 square feet, that's a square meter of paint per \$300 of real estate. I vaguely remember that the world's real estate has a combined value of something like \$100 trillion, [4] which suggests there's about 300 billion square meters of paint on the world's real estate. That's about one New Mexico.

NOT ENOUGH	EXACTLY ENOUGH	MORE THAN ENOUGH
//	/	

Of course, both of the building-related guesses could be overestimates (lots of buildings are not painted) or underestimates (lots of things that are not buildings [5] are painted) But from these wild Fermi estimates, my guess would be that there probably isn't enough paint to cover all the land.

So, how did Fermi do?

According to the report [**The State of the Global Coatings Industry**](#), the world produced 34 billion liters of paints and coatings in 2012.

There's a neat trick that can help us here. If some quantity—say, the world economy—has been growing for a while at an annual rate of n —say, 3% (0.03)—then the most recent year's share of the whole total so far is $1 - \frac{1}{1+n}$, and the whole total so far is the most recent year's amount times $1 + \frac{1}{n}$.

If we assume paint production has, in recent decades, followed the economy and grown at about 3% per year, that means the total amount of paint produced equals the current yearly production times 34.^[6] That comes out to a little over a trillion liters of paint. At 30 square meters per gallon,^[7] that's enough to cover 9 trillion square meters—about the area of the United States.

So the answer is no; there's not enough paint to cover the Earth's land, and—at this rate—probably won't be enough until the year 2100.

Data Intuition

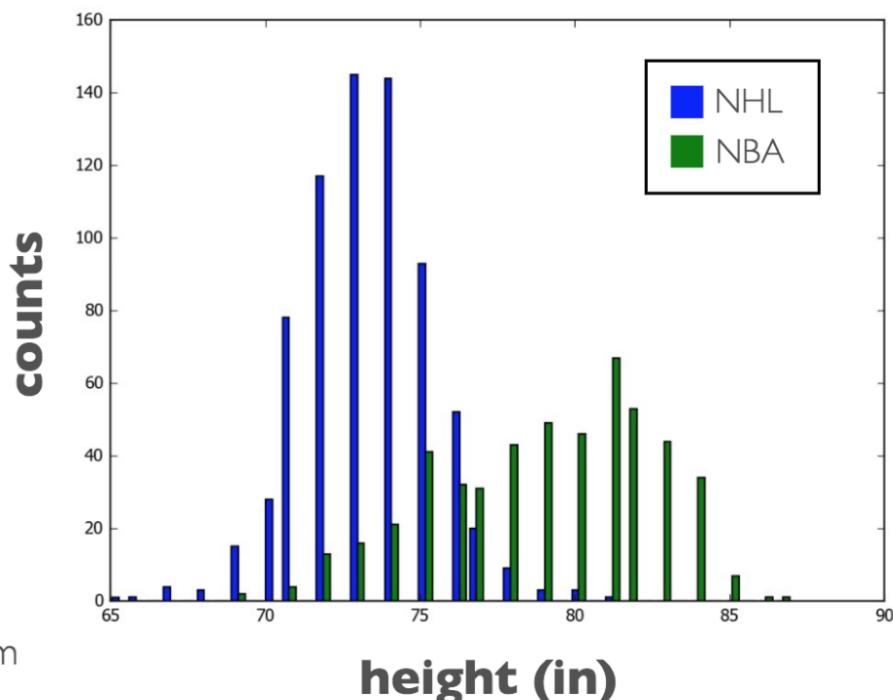
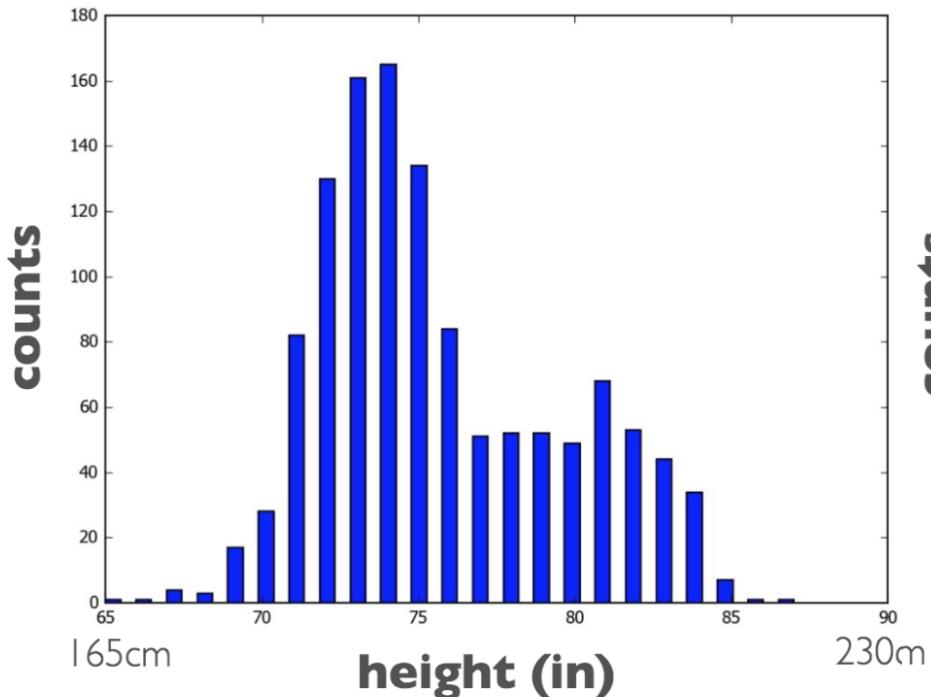
1. Think about your question and your expectations
2. Do some Fermi calculations (back of the envelope calculations)
3. Write code & look at outputs <- think about those outputs
4. Use your gut instinct / background knowledge to guide you
5. Review code & fix bugs

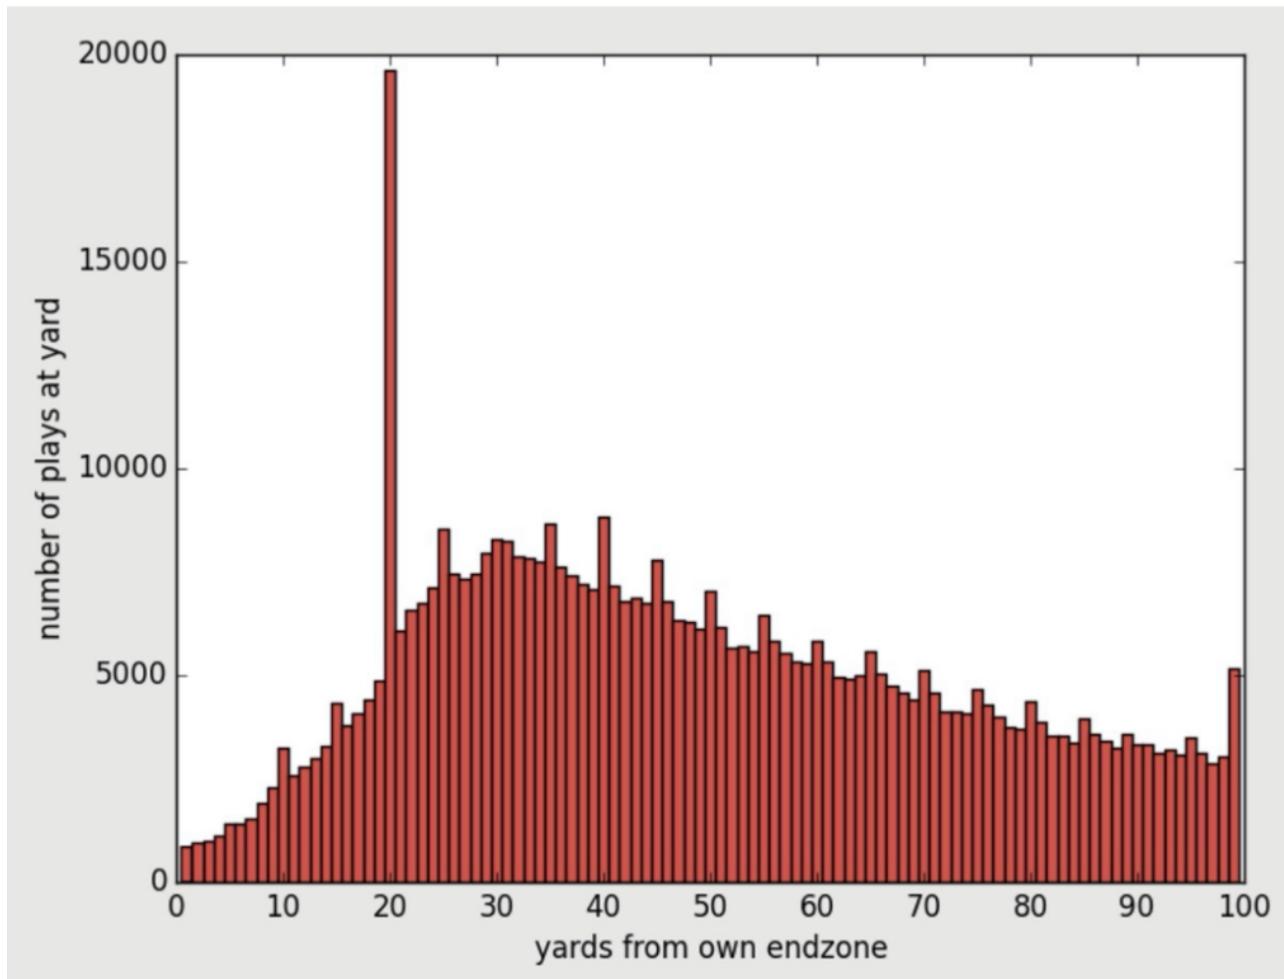
On your own (meaning w/o Googling), please fill out quickly:

<https://forms.gle/CREcpMkYDLYTUp2s6>



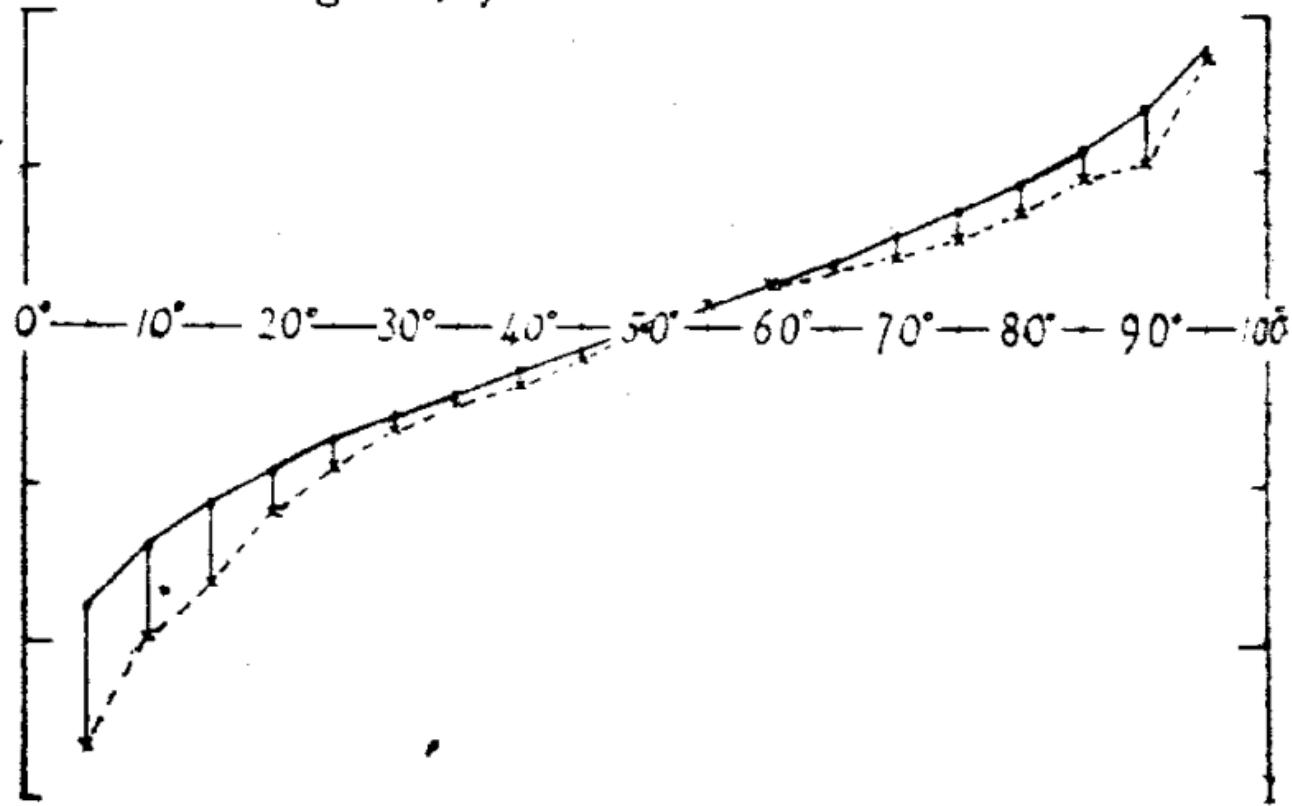
Other kinds of
guessing





Diagram, from the tabular values.

Vox Populi



The Wisdom of the Crowds

- Diversity of opinion: Each person should have private information....even if it's just an eccentric interpretation of the known facts
- Independence: People's opinions aren't determined by the opinions of those around them
- Decentralization: People are able to specialize and draw on local knowledge
- Aggregation: Some mechanism exists for turning private judgements into a collective decision